Quantum Computation as Geometry

Michael A. Nielsen,* Mark R. Dowling, Mile Gu, Andrew C. Doherty

Quantum computers hold great promise for solving interesting computational problems, but it remains a challenge to find efficient quantum circuits that can perform these complicated tasks. Here we show that finding optimal quantum circuits is essentially equivalent to finding the shortest path between two points in a certain curved geometry. By recasting the problem of finding quantum circuits as a geometric problem, we open up the possibility of using the mathematical techniques of Riemannian geometry to suggest new quantum algorithms or to prove limitations on the power of quantum computers.

Quantum computers have the potential to solve efficiently some problems that are considered intractable on conventional classical computers: The most famous example is Shor's algorithm (*1*) for finding the prime factors of an integer. Despite this great promise, as yet there is no general method for constructing good quantum algorithms, and very little is known about the potential power (or limitations) of quantum computers.

A quantum computation is usually described as a sequence of logical gates, each coupling only a small number of qubits. The sequence of gates determines a unitary evolution U performed by the computer. The difficulty of performing the computation is characterized by the number of gates used by the algorithm, which is said to be efficient if the number of gates required grows only polynomially with the size of problem (e.g., with the number of digits in the number to be factored, in the case of Shor's factoring algorithm).

We developed an alternate approach to understanding the difficulty of implementing a unitary operation U. Suppose that U is generated by some time-dependent Hamiltonian $H(t)$ according to the Schrödinger equation $dU/dt = -iHU$, where i is $\sqrt{-1}$ and with the requirement that at an appropriate final time t_f, $U(t_f) = U$. We characterized the difficulty of the computation by imposing a cost $F[H(t)]$ on the Hamiltonian control, $H(t)$. Following (*2*), we chose a cost function on $H(t)$ that defines a Riemannian geometry on the space of unitary operations. Finding the optimal control function $H(t)$ for synthesizing a desired unitary U then corresponds to finding minimal geodesics of the Riemannian geometry.

We show here that the minimal geodesic distance between the identity operation I and U is essentially equivalent to the number of gates required to synthesize U. This result extends the work in (*2*), where it was shown that the minimal distance provides a lower bound on the number of gates required to synthesize U.

Our result allows the tools of Riemannian geometry to be applied to understand quantum computation. In particular, we can use a powerful tool—the calculus of variations—to find the geodesics of the space. Just as in general relativity, this calculus can be used to derive the geodesic equation, a "force law" whereby the local shape of space tells us how to move in order to follow the geodesics of the manifold.

Intuitively, our results show that the optimal way of solving any computational problem is to "fall freely" along the minimal geodesic curve connecting the identity operation to the desired operation, with the motion determined entirely by the local "shape" of the space. To appreciate this result, consider that once an initial position and velocity are set, the remainder of the geodesic is completely determined by the geodesic equation. This is in contrast with the usual case in circuit design, either classical or quantum, where being given part of an optimal circuit does not obviously assist in the design of the rest of the circuit. Geodesic analysis thus offers a potentially powerful approach to the analysis of quantum computation. However, a caveat to this optimism is that although we know the initial position is the identity operation, we still need to determine the initial velocity in order to find the minimal geodesic; this is not, in general, an easy problem.

Our results can also be viewed as showing that the problem of finding minimal quantum circuits is equivalent to a problem in geometric control theory (*3*), which has had great success in using techniques from the calculus of variations and Riemannian geometry to solve optimal control problems. For example, Khaneja and co-workers (*4*) and others (*5*, *6*) have used geometric techniques to analyze the minimal time cost of synthesizing two-qubit unitary operations using a fixed, two-qubit control Hamiltonian and fast local control.

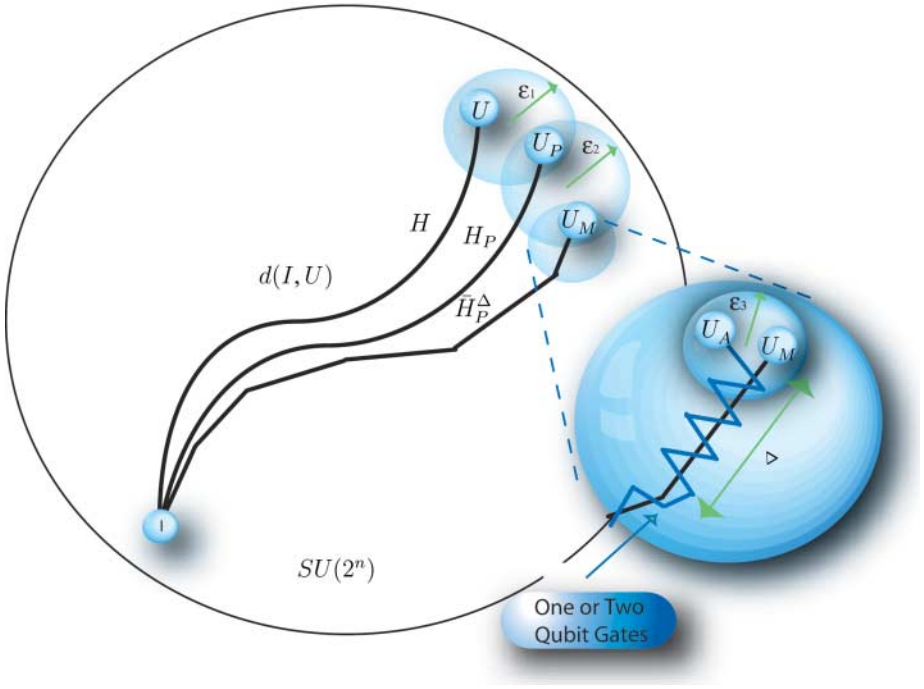

Fig. 1. Schematic of the three steps used to construct a quantum circuit approximating the unitary operation U. The circuit is of a size that is polynomial in the distance $d(I, U)$ between the identity and U. First, we projected the Hamiltonian $H(t)$ for the minimal geodesic path onto one- and two-qubit terms, giving $H_P(t)$. By choosing the penalty p large enough ($p = 4^n$), we ensured the error in this approximation is small, $\varepsilon_1 \leq d(I, U)/2^n$. Next, we broke up the evolution according to $H_P(t)$ into N small time steps of size $\Delta = d(I, U)/N$, and we approximated with a constant mean Hamiltonian $\bar{H}_P^{\Delta}$ over each step. Finally, we approximated evolution according to the constant mean Hamiltonian over each step by a sequence of one- and two-qubit quantum gates. The total errors, ε_2 and ε_3, introduced by these approximations can be made smaller than any desired constant by choosing the step size Δ sufficiently small: $\Delta = O(1/[n^2 d(I, U)])$. In total, we need $O(n^6 d(I, U)^3)$ quantum gates to approximate U to within some constant error, which can be made arbitrarily small.

School of Physical Sciences, The University of Queensland, Queensland 4072, Australia.

*To whom correspondence should be addressed. E-mail: nielsen@physics.uq.edu.au

To choose a cost function on the control Hamiltonian $H(t)$, we first write $H(t)$ in terms of the Pauli operator expansion $H = \sum_{\sigma}' h_\sigma \sigma + \sum_{\sigma}'' h_\sigma \sigma$, assuming the following: (i) In the first sum, σ ranges over all possible one- and two-body interactions, that is, over all products of either one or two Pauli matrices acting on n qubits. (ii) In the second sum, σ ranges over all other tensor products of Pauli matrices and the identity. (iii) The h_σ are real coefficients. We then define a measure of the cost of applying a particular Hamiltonian during synthesis of a desired unitary operation

$$F(H) = \sqrt{\sum_{\sigma}' h_\sigma^2 + p^2 \sum_{\sigma}'' h_\sigma^2} \qquad (1)$$

The parameter p is a penalty paid for applying three- and more-body terms; later we will choose p to be large in order to suppress such terms (7).

This definition of control cost leads us to a natural notion of distance in the space $SU(2^n)$ of n-qubit unitary operators with unit determinant. A curve $[U]$ between the identity operation I and the desired operation U is a smooth function U: $[0, t_f] \rightarrow SU(2^n)$, such that $U(0) = I$ and $U(t_f) = U$. The length of this curve can then be defined by the total cost of synthesizing the Hamiltonian that generates evolution along the curve

$$d([U]) \equiv \int_0^{t_f} dt F[H(t)] \qquad (2)$$

Because $d([U])$ is invariant with respect to different parameterizations of $[U]$ (*8*), we can always rescale the Hamiltonian $H(t)$ such that $F[H(t)] = 1$ and the desired unitary U is generated at time $t_f = d([U])$. From now on, we assume that we are working with such normalized curves. Finally, the distance $d(I, U)$ between I and U is defined to be the minimum of $d([U])$ over all curves $[U]$ connecting I and U.

We will show that for any family of unitaries U (implicitly, U is indexed by the number of qubits n), there is a quantum circuit containing a number of gates that is polynomial in $d(I, U)$ and that approximates U to high accuracy. In other words, if the distance $d(I, U)$ scales polynomially with n for some family of unitary operations, then it is possible to find a polynomial-sized quantum circuit for that family of unitary operations. Conversely, the metric we construct also has the property, proved in (*2*), that up to a constant factor, the distance $d(I, U)$ is a lower bound on the number of one- and two-qubit quantum gates required to exactly synthesize U. Consequently, the distance $d(I, U)$ is a good measure of the difficulty of implementing the operation U on a quantum computer.

The function $F(H)$ specified by eq. 1 can be thought of as the norm associated to a (right invariant) Riemannian metric whose metric tensor g has components:

$$g_{\sigma\tau} = \begin{cases} 0 & \text{if } \sigma \neq \tau \\ 1 & \text{if } \sigma = \tau \text{ and } \sigma \text{ is one- or two-body} \\ p^2 & \text{if } \sigma = \tau \text{ and } \sigma \text{ is three- or more-body} \end{cases} \qquad (3)$$

These components are written with respect to a basis for the local tangent space corresponding to the Pauli expansion coefficients h_σ. The distance $d(I, U)$ is equal to the minimal length solution to the geodesic equation, which may be written (*9*) as $\langle dH/dt, K\rangle = i\langle H, [H, K]\rangle$. In this expression, the notation $\langle x, y\rangle$ indicates the inner product of x and y on the tangent space $su(2^n)$ defined by the metric components of eq. 3, the notation $[H, K]$ indicates the matrix commutator, and K is an arbitrary operator. For our particular choice of metric components, this geodesic equation may be rewritten as

$$p_\sigma^2 \dot{h}_\sigma = i \sum_\tau p_\tau^2 h_\tau \tilde{h}_{[\sigma,\tau]} \qquad (4)$$

where $\tilde{h}_{[\sigma,\tau]} = \mathrm{tr}(H[\sigma, \tau])/2^n$ and tr indicates the trace. A particular class of solutions to this equation was studied in (*2*), but understanding the general behavior of the geodesics remains a problem for future research (*10*). There are powerful tools in Riemannian geometry (*11*, *12*) available for the study of minimal length geodesics.

Our goal is to use the optimal control Hamiltonian $H(t)$ to construct explicitly a quantum circuit containing a number of gates that is polynomial in $d(I, U)$ and which approximates U closely. The construction combines three main ideas, which we express through three separate lemmas, before combining them to obtain the result (Fig. 1).

The first lemma shows that the error that arises by simply ignoring the many-body interactions in $H(t)$ can be made small by choosing the penalty p appropriately. We define H_P to be the projected Hamiltonian formed by deleting all three- and more-body terms in the Pauli expansion. The following result is proved in (*13*).

Lemma 1: Let $H_P(t)$ be the projected Hamiltonian obtained from a Hamiltonian $H(t)$ generating a unitary U. Let U_P be the corresponding unitary generated by $H_P(t)$. Then

$$\|U - U_P\| \leq \frac{2^n d([U])}{p} \qquad (5)$$

where $\|x\|$ is the operator norm of x (*14*) and p is the penalty parameter appearing in the definition of the metric. Thus, by choosing p sufficiently large, say $p = 4^n$, we can ensure that $\|U - U_P\| \leq d([U])/2^n$.

Motivated by the preceding lemma, we change our aim from accurately synthesizing U to accurately synthesizing U_P. To do this, we break the evolution according to $H_P(t)$ up into many small intervals, each of length Δ. The next lemma shows that evolution according to the time-dependent Hamiltonian $H_P(t)$ over such a small time interval can always be accurately simulated by a constant mean Hamiltonian, which we denote $\overline{H}_P^\Delta$.

Lemma 2: Let U be an n-qubit unitary generated by applying a time-dependent Hamiltonian $H(t)$ satisfying $\|H(t)\| \leq c$, for some constant c, over a time interval $[0, \Delta]$. Then defining the mean Hamiltonian $\overline{H} \equiv \frac{1}{\Delta}\int_0^\Delta dt H(t)$ we have

$$\|U - \exp(-i\overline{H}\Delta)\| \leq 2(e^{c\Delta} - 1 - c\Delta) = O(c^2\Delta^2) \qquad (6)$$

where $O(x)$ indicates the asymptotic behavior of the function. The proof of this lemma is based on the Dyson operator expansion and is presented in (*13*). To apply this lemma to $H_P(t)$, note that elementary norm inequalities and the observation $F[H_P(t)] \leq 1$ imply that $\|H_P(t)\| \leq (3/\sqrt{2})nF[H_P(t)] \leq (3/\sqrt{2})n$ (*15*). Lemma 2 implies that over a time interval Δ, we have

$$\|U_P^\Delta - \exp(-i\overline{H}_P^\Delta\Delta)\| \leq 2\left[e^{3/\sqrt{2}n\Delta} - (1 + \frac{3}{\sqrt{2}}n\Delta)\right] = O(n^2\Delta^2) \qquad (7)$$

where U_P^Δ is the evolution generated by $H_P(t)$ over the time interval Δ, and $\overline{H}_P^\Delta$ is the corresponding mean Hamiltonian.

Our third and final lemma shows that evolution according to a time-independent Hamiltonian H containing only one- and two-body terms can be very accurately simulated by using a number of quantum gates that is not too large.

Lemma 3: Suppose H is *an* n-qubit two-body Hamiltonian whose Pauli expansion coefficients satisfy $|h_\sigma| \leq 1$. Then there exists a unitary U_A, satisfying

$$\|e^{-iH\Delta} - U_A\| \leq c_2 n^4 \Delta^3 \qquad (8)$$

that can be synthesized using at most $c_1 n^2/\Delta$ one- and two-qubit gates, where c_1 and c_2 are constants.

This result follows from standard procedures for simulating quantum evolutions using quantum gates [(*16*) chap. 4], and it is proved in (*13*). The average Hamiltonian $\overline{H}_P^\Delta$ provided by lemma 2 satisfies the assumptions of lemma 3, because the Pauli expansion coefficients of $H_P(t)$ satisfy $|h_\sigma| \leq 1$ for all times.

To integrate lemmas 1 to 3, suppose $H(t)$ is the time-dependent normalized Hamiltonian generating the minimal geodesic of length $d(I, U)$. Let $H_P(t)$ be the corresponding projected

Hamiltonian, which generates U_P and satisfies $\|U - U_P\| \leq d(I, U)/2^n$, as guaranteed by lemma 1, and where we have chosen $p = 4^n$ as the penalty. Now divide the time interval $[0, d(I, U)]$ up into a large number N of time intervals each of length $\Delta = d(I, U)/N$. Let U_P^j be the unitary operation generated by $H_P(t)$ over the jth time interval, where j is an integer. Let U_M^j be the unitary operation generated by the corresponding mean Hamiltonian. Then lemma 2 implies that:

$$\|U_P^j - U_M^j\| \leq 2[e^{3\sqrt{2}n\Delta} - (1 + \frac{3}{\sqrt{2}}n\Delta)] \quad (9)$$

Lemma 3 implies that we can synthesize a unitary operation U_A^j using at most $c_1 n^2/\Delta$ one- and two-qubit gates and satisfying $\|U_M^j - U_A^j\| \leq c_2 n^4 \Delta^3$.

Putting all these results together and applying the triangle inequality repeatedly, we obtain

$$\|U - U_A\| \leq \|U - U_P\| + \|U_P - U_A\| \quad (10)$$

$$\leq \frac{d(I,U)}{2^n} + \sum_{j=1}^{N} \|U_P^j - U_A^j\| \quad (11)$$

$$\leq \frac{d(I,U)}{2^n} + \sum_{j=1}^{N} \left(\|U_P^j - U_M^j\| + \|U_M^j - U_A^j\| \right) \quad (12)$$

$$\leq \frac{d(I,U)}{2^n} + 2\frac{d(I,U)}{\Delta} \times \left[e^{(3/\sqrt{2})n\Delta} - \left(1 + \frac{3}{\sqrt{2}}n\Delta\right) \right] + c_2 d(I,U) n^4 \Delta^2 \quad (13)$$

Provided we choose Δ to scale at most as $1/[n^2 d(I, U)]$, we can ensure that the error in our approximation U_A to U is small, and the number of gates scales as $n^6 d(I, U)^3$.

Summing up, we have the following theorem (*17*): Using $O(n^6 d(I, U)^3)$ one- and two-qubit gates, it is possible to synthesize a unitary U_A satisfying $\|U - U_A\| \leq c$, where c is any constant (e.g., $c = 1/10$).

Our results demonstrate that, up to polynomial factors, the optimal way of generating a unitary operation is to move along the minimal geodesic curve connecting I and U. Because the length of such geodesics also provides a lower bound on the minimal number of quantum gates required to generate U, as shown in (*2*), the geometric formulation offers an alternate approach, which may suggest efficient quantum algorithms or provide a way of proving that a given algorithm is indeed optimal.

It would, of course, be desirable to completely classify the geodesics of the metric we constructed. An infinite class of such geodesics has been constructed in (*2*) and is shown to have an intriguing connection to the problem of finding the closest vector in a lattice. A more complete classification of the geodesics could provide major insight on the potential power of quantum computation.

References and Notes

1. P. W. Shor, *Proceedings of the 35th Annual Symposium on Fundamentals of Computer Science* (IEEE Press, Los Alamitos, CA, 1994).
2. M. A. Nielsen, *Quant. Inf. Comput.*, in press (preprint available at http://arxiv.org/abs/quant-ph/0502070).
3. V. Jurdjevic, *Geometric Control Theory* (Cambridge Univ. Press, Cambridge, 1996).
4. N. Khaneja, R. Brockett, S. J. Glaser, *Phys. Rev. A* **63**, 032308 (2001).
5. N. Khaneja, S. J. Glaser, *Chem. Phys.* **267**, 11 (2001).
6. N. Khaneja, S. J. Glaser, R. Brockett, *Phys. Rev. A* **65**, 032301 (2002).
7. An alternate way of viewing this cost function is as a penalty metric of the kind used in sub-Riemannian geometry (*18*, p. 18).
8. The cost function has the property that $F(\alpha H) = |\alpha| F(H)$, where α is any real number. This property and the chain rule imply invariance of the length with respect to reparameterization.
9. V. I. Arnold, B. A. Khesin, *Topological Methods in Hydrodynamics*, vol. 125 of *Applied Mathematical Sciences* (Springer, New York, 1998).
10. Our metric is superficially similar to the usual metric of Euclidean space, and it is tempting to suppose that geodesics must be straight lines, i.e., constant Hamiltonians. However, the Pauli coefficients for the Hamiltonian actually correspond to a changing local basis for the tangent space, not a fixed basis, and hence constant Hamiltonians do not, in general, give rise to geodesics.
11. J. Milnor, *Morse Theory* (Princeton Univ. Press, Princeton, NJ, 1969).
12. M. Berger, *A Panoramic View of Riemannian Geometry* (Springer, Berlin, 2003).
13. Proofs are available as supporting material on *Science* Online.
14. The operator norm of X is defined as $\|X\| = \max_{|\psi\rangle} |\langle\psi|X|\psi\rangle|$, where the maximization is over all normalized vectors, $|\langle\psi|\psi\rangle|^2 = 1$.
15. The first inequality comes from the fact that there are $9n(n - 1)/2 + 3n$ one- and two-qubit terms.
16. M. A. Nielsen, I. L. Chuang, *Quantum Computation and Quantum Information* (Cambridge Univ. Press, Cambridge, 2000).
17. The overhead factors in this theorem may be substantially improved, e.g., by making use of higher order analyses in lemmas 1 to 3. However, the key point—that U can be accurately approximated with a number of gates that scales polynomially with $d(I, U)$—remains the same.
18. R. Montgomery, *A Tour of Subriemannian Geometries, Their Geodesics and Applications*, vol. 91 of *Mathematical Surveys and Monographs* (American Mathematical Society, Providence, RI, 2002).

19. We thank S. Aaronson, M. de Burgh, J. Dodd, C. Hill, A. Hines, A. Lund, L. Noakes, M. Sarovar, and B. Toner for helpful discussions and the Australian Research Council for funding. We are especially grateful to L. Noakes for pointing out to us the simplified form of the geodesic equation for right-invariant metrics.

Supporting Online Material
www.sciencemag.org/cgi/content/full/311/5764/1133/DC1
SOM Text

19 October 2005; accepted 20 December 2005
10.1126/science.1121541

Effects of Solar Flares on the Ionosphere of Mars

Michael Mendillo,[1] Paul Withers,[1]* David Hinson,[2] Henry Rishbeth,[1,3] Bodo Reinisch[4]

All planetary atmospheres respond to the enhanced x-rays and ultraviolet (UV) light emitted from the Sun during a flare. Yet only on Earth are observations so continuous that the consequences of these essentially unpredictable events can be measured reliably. Here, we report observations of solar flares, causing up to 200% enhancements to the ionosphere of Mars, as recorded by the Mars Global Surveyor in April 2001. Modeling the altitude dependence of these effects requires that relative enhancements in the soft x-ray fluxes far exceed those in the UV.

Sudden changes in the Sun's photon radiation and in the particles and fields of its solar wind reach Earth in about 8 min and a few days, respectively. These enhanced sources of energy cause sudden atmospheric disturbances and the auroral displays associated with longer lived geomagnetic storms. The recent availability of spacecraft orbiting other planets has enabled studies of such effects on other worlds. A mass ejection from the Sun's corona in early November 2000 caused auroras on Earth, Jupiter, and Saturn during its month-long traverse through the solar system, providing a specific challenge to models that track solar wind density and magnetic field enhancements (*1*). Increased x-ray emissions were observed from Jupiter and Saturn in November 2003 and January 2004, respectively, shortly after solar flares, thereby demonstrating the Sun's control of nonauroral x-ray emission from giant planets (*2*, *3*). However, the direct response of another planetary atmosphere to solar flare photons, e.g., suddenly enhancing its ionosphere, has not been seen. Here, we report such an effect in the ionosphere of Mars.

Ions and electrons in a planet's ionosphere are produced by the photoionization of neutral

[1]Center for Space Physics, Boston University, Boston, MA 02215, USA. [2]Department of Electrical Engineering, Stanford University, Stanford, CA 94305, USA. [3]School of Physics and Astronomy, University of Southampton, Southampton SO17 1BJ, UK. [4]Center for Atmospheric Research, University of Massachusetts Lowell, Lowell, MA 01854, USA.

*To whom correspondence should be addressed. E-mail: withers@bu.edu